

Operational Analytics - project-3

1. Project Description

This project focuses on **Operational Analytics**, aiming to investigate company-wide operations using SQL. The analysis involves two case studies:

1. **Job Data Analysis** – Understanding performance and quality across job reviews.
2. **Metric Spike Investigation** – Analyzing behavioral trends and engagement metrics to detect changes in user activity.

2. Project Approach

- Data was provided in `.csv` files and imported into Pandas (Python) for exploratory analysis and SQL logic simulation.
- Each business question was broken into tasks, followed by SQL query simulation using DataFrame operations.
- Resulting insights were interpreted and summarized to guide strategic recommendations.

3. Tech Stack

- **Python (Pandas)**: For simulating SQL queries and data manipulation.
- **MySQL Workbench**: Intended SQL engine for final deployment.
- **Excel/Google Sheets**: For exporting output snapshots.
- **Google Drive**: For storing and sharing final report.

4. Case Study 1: Job Data Analysis

Task 1: Jobs Reviewed Per Hour (Nov 2020)

SQL Query:

```
SELECT
  DATE(ds) AS review_date,
  HOUR(TIMESTAMP(ds)) AS review_hour,
  COUNT(*) AS jobs_reviewed
FROM job_data
WHERE ds BETWEEN '2020-11-01' AND '2020-11-30'
GROUP BY review_date, review_hour
ORDER BY review_date, review_hour;
```

Output :

ds	hour	jobs_reviewed
2020-11-25	18	1
2020-11-26	6	1
2020-11-27	20	1

Conclusion:

Hourly review activity was generally balanced with low concurrency. Useful for detecting peak workflow hours.

Insights:

- No visible clustering of job reviews; possibly auto-assigned or evenly distributed workflows.
- Recommending more structured time-window allocation for reviewer efficiency.

Task 2: 7-Day Rolling Throughput

SQL Query:

```
WITH daily_stats AS (  
    SELECT  
        ds,  
        COUNT(*) AS total_jobs,  
        SUM(time_spent) AS total_time  
    FROM job_data  
    GROUP BY ds  
)  
SELECT  
    ds,  
    total_jobs / total_time AS throughput,  
    AVG(total_jobs / total_time) OVER (ORDER BY ds ROWS BETWEEN 6  
PRECEDING AND CURRENT ROW) AS rolling_avg_7d  
FROM daily_stats;
```

Output :

ds	throughput	rolling_avg_7d
2020-11-28	0.0606	NULL
2020-11-29	0.0500	NULL

Conclusion:

Rolling averages smooth out anomalies, offering better trend analysis over time.

Insights:

- Sudden throughput spikes could indicate system errors or batch task releases.

- Recommend monitoring both raw and smoothed metrics for balance.

Task 3: Language Share (Last 30 Days)

SQL Query:

```
SELECT
  language,
  ROUND(100.0 * COUNT(*) / (SELECT COUNT(*) FROM job_data
    WHERE ds >= DATE_SUB(MAX(ds), INTERVAL 30 DAY)), 2) AS share
FROM job_data
WHERE ds >= DATE_SUB(MAX(ds), INTERVAL 30 DAY)
GROUP BY language;
```

Output:

Language	Share (%)
Persian	37.5
English	12.5
Arabic	12.5
Hindi	12.5
French	12.5

Insights:

- Persian dominates the review language pool. May indicate geographical focus.
- Recommending localization improvements and language-based staffing allocation.

Task 4: Duplicate Detection

SQL Query:

```
SELECT *  
FROM job_data  
GROUP BY job_id, actor_id, event, language, time_spent, org, ds  
HAVING COUNT(*) > 1;
```

Output:

- **No duplicate rows detected.**

Insights:

- Data integrity confirmed. No remedial action needed for duplicates.

Case Study 2: Investigating Metric Spike

Task 1: Weekly User Engagement

SQL Query:

```
SELECT
  DATE_TRUNC('week', occurred_at) AS week,
  COUNT(DISTINCT user_id) AS active_users
FROM events
GROUP BY week;
```

Output :

Week	Active Users
2014-04-28	701
2014-05-05	1054

Insights:

- Consistent weekly increase in engagement; useful for retention analysis and campaign planning.

Task 2: User Growth Over Time

SQL Query:

```
SELECT
  DATE_TRUNC('week', created_at) AS week,
  COUNT(user_id) OVER (ORDER BY DATE_TRUNC('week', created_at))
  AS cumulative_users
FROM users;
```

Output :

Week	Cumulative Users
2013-01-07	55
2013-01-14	102

Insights:

- Growth trajectory is healthy and incremental, signaling good acquisition strategy.

Task 3: Weekly Retention by Cohort

SQL Query:

```
SELECT
  signup_week,
  event_week,
  COUNT(DISTINCT user_id) AS retained_users
FROM (
  SELECT
    u.user_id,
    DATE_TRUNC('week', u.created_at) AS signup_week,
    DATE_TRUNC('week', e.occurred_at) AS event_week
  FROM users u
  JOIN events e ON u.user_id = e.user_id
) t
GROUP BY signup_week, event_week;
```

Output :

Signup Week	Event Week	Retained Users
2012-12-31	2014-04-28	2

Insights:

- Small but consistent retention across long periods. Indicates long-term value from early users.

Task 4: Weekly Engagement Per Device

SQL Query:

```
SELECT
  DATE_TRUNC('week', occurred_at) AS week,
  device,
  COUNT(DISTINCT user_id) AS active_users
FROM events
GROUP BY week, device;
```

Output :

Week	Device	Active Users
2014-04-28	ASUS Chromebook	23

Insights:

- Chromebook and mobile engagement are higher than desktop in early weeks. Recommend responsive UI focus.

Task 5: Email Engagement

SQL Query:

```
SELECT
  DATE_TRUNC('week', occurred_at) AS week,
  action,
  COUNT(*) AS count
FROM email_events
GROUP BY week, action;
```

Output :

Week	Action	Count
2014-04-28	email_clickthrough	187
2014-04-28	sent_weekly_digest	908

Insights:

- Digest emails are the most frequent and should be optimized for higher click-through rates.
- Re-engagement campaigns need follow-up tracking.